

# Eco GenStage

## Exercises

### Setup

In this section of the exercises, the working directory is `eco` folder within the `beam2024` directory.

The project was generated as an umbrella application. Feel free to explore folder structure and file contents. The command used for project generation was:

```
$ mix phx.new eco --no-mailer --no-ecto --umbrella --no-gettext
```

A key point to note is that this project was created as a Phoenix application rather than a pure Elixir application. The command accepts the same parameters as `mix new`, along with additional options for managing components such as Gettext for internationalization or mailer functionality.

To ensure you have Hex and Phoenix, execute the following commands:

```
$ mix local.hex  
$ mix archive.install hex phx_new
```

Next, set up the application (using a predefined alias in the `mix.exs` file):

```
$ mix setup
```

To verify that everything is running smoothly run the below command and navigate to `localhost:4000` in the browser:

```
$ iex -S mix phx.server
```

### Goals

The application revolves around submitting sample data to process it with GenStage and finally assigning a star to a respective bin depending on the string contents.

### Data processing pipeline

**Inputs** The application expects data to be submitted via a form, located in the template in the `index_live.ex`. The form contains only one field and a submit button. The data is accepted as a string of any length and is case-insensitive.

**Outputs** The pipeline should publish an event on the `recyclables` topic with a `{:recycled, sign}` payload using Phoenix.PubSub.

**Producing events - Producer / Source** In the `apps/eco/lib/eco/` folder, create a `source.ex` file. This will be the GenStage producer. The general structure should be something like this:

```

defmodule Eco.Source do
  use GenStage

  def start_link(initial_state), do GenStage.start_link(__MODULE__, initial_state, name: __MODULE__)

  def init(initial_state) do
    # your code here
  end

  def handle_demand(demand, state) do
    # your code here
  end
end

```

To correctly initialize the producer, return `{:producer, initial_state, dispatcher: GenStage.BroadcastDispatcher}` from the `init/1` function.

Since we want to leverage the abilities of GenStage to dispatch events automatically as they come in, we don't need a complex `handle_demand/2` implementation. Set the return value to `{:noreply, [], state}`.

To dispatch events down the pipeline, using the form input from the user, we will add `handle_call/3`. Add the following functions:

```

def start(event, timeout \\ 5000) do
  GenStage.call(__MODULE__, {:start, sequence}, timeout)
end

def handle_call({:start, sequence}, _from, state) do
  events = [],
  {:reply, :ok, events, state}
end

```

As mentioned in the Inputs section, our data at this point should be the submitted string. To simulate batch processing, split the string into a list of single letters.

For example, if the input is "ACTGMGPBNHDACT", the output value for the events variable should be ["A", "C", "T", ...]. Note that at this stage, we are only generating events; no additional data processing is taking place here.

**Event data processing** In the `apps/eco/lib/eco/` directory, create a `producer_consumer.ex` file. This will correspond to the GenStage `producer_consumer`, a stage where our data will undergo transformation. Here, the data will be converted to upper case and some values will be discarded. The general structure should be something like this:

```

defmodule Eco.ProducerConsumer do
  use GenStage

```

```

def start_link(_initial) do
  GenStage.start_link(__MODULE__, [], name: __MODULE__)
end

@impl true
def init(init_state), do: {:producer_consumer, init_state, subscribe_to: [Eco.Source]}

@impl true
def handle_events(events, _from, state) do
  processed_events = []
  {:noreply, processed_events, state}
end
end

```

The `start_link/1` is implemented to allow us to incorporate the `ProducerConsumer` module into the supervision tree.

The `init/1` function marks the process as a `:producer_consumer` and subscribes it to the `Eco.Source` module.

In `handle_events/3`, implement the event processing logic - these events will be published to the consumer. Since the event content is case-insensitive, we can freely change its case to facilitate comparison with exemplary data. Apply `String.upcase/1` to the event content and discard any events whose content is not included in the list `["A", "C", "T", "G"]`.

**Consume events** In the `apps/eco/lib/eco/` directory, create a `sink.ex` file. This module will consume events produced by the previous stages. Here's the general module implementation:

```

defmodule Eco.Sink do
  use GenStage
  alias Phoenix.PubSub

  def start_link(_initial) do
    GenStage.start_link(__MODULE__, [])
  end

  def init(state) do
    {:consumer, state, subscribe_to: [Eco.ProducerConsumer]}
  end

  def handle_events(events, _from, state) do
    # your code here
    {:noreply, [], state}
  end
end

```

`end`

The `start_link/1` is implemented to allow us to incorporate the `Sink` module into the supervision tree.

The `init/1` function marks the process as a `:consumer` and subscribes it to the `Eco.ProducerConsumer` module.

In `handle_events/3`, implement the logic to consume the list of events. At this stage, iterate over the `events` variable (similar to what we did in the `ProducerConsumer` module) and publish messages on a specific topic using `PubSub`.

To publish messages, use `Phoenix.PubSub.broadcast/3`. The topic to publish information to is `"recyclables"` and the message format should be `{:recycled, sign}`, where `sign` is a value from the event.

**Tying it all up** Before our pipeline can accept the request, we need to include it in our application's children. To do so, modify the list of children in the `Eco.Application` module by adding the following:

```
children = [  
  # ...  
  {Eco.Source, []},  
  {Eco.ProducerConsumer, []},  
  {Eco.Sink, []}  
  # ...  
]
```

Lastly, in `index_live.ex`, modify the form submit event (remember to add an alias to `Source`):

```
def handle_event("save", %{"sequence" => sequence}, socket) do  
  Source.start(sequence)  
  {:noreply, socket}  
end
```

**Additional** Start multiple `Sink` instances, modify `handle_event/3` to sleep for a random time between 0.5s and 2s with `Process.sleep(Enum.random(500..2000))`, and observe what happens.